

```
In [2]: import pandas as pd
import numpy as np
```

```
In [3]: # Reading the new CSV file
new_data = pd.read_csv('last_two_years_accidents.csv')
```

```
In [32]: new_data.shape
```

```
Out[32]: (2009085, 42)
```

```
In [33]: # Checking for missing values
missing_values = new_data.isnull().sum()
print(missing_values)
```

ID	0
Source	0
Severity	0
Start_Time	0
End_Time	0
Start_Lat	0
Start_Lng	0
Distance(mi)	0
Street	0
City	0
County	0
State	0
Zipcode	0
Country	0
Timezone	0
Airport_Code	0
Weather_Timestamp	0
Temperature(F)	0
Humidity(%)	0
Pressure(in)	0
Visibility(mi)	0
Wind_Direction	0
Wind_Speed(mph)	0
Weather_Condition	0
Amenity	0
Bump	0
Crossing	0
Give_Way	0
Junction	0
No_Exit	0
Railway	0
Roundabout	0
Station	0
Stop	0
Traffic_Calming	0
Traffic_Signal	0
Turning_Loop	0
Sunrise_Sunset	0
Civil_Twilight	0
Nautical_Twilight	0
Astronomical_Twilight	0
Cluster	0

dtype: int64


```
In [34]: from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [35]: # Selecting relevant features for clustering (numeric columns)
numeric_features = new_data.select_dtypes(include=['float64', 'int64']).columns
new_data_numeric = new_data[numeric_features]
```

```
In [36]: # Standardizing the numeric features
scaler = StandardScaler()
new_data_scaled = scaler.fit_transform(new_data_numeric)
```

```
In [37]: # Applying K-means clustering with 3 clusters
kmeans = KMeans(n_clusters=3, random_state=42)
new_data['Cluster'] = kmeans.fit_predict(new_data_scaled)
```

```
In [38]: # Visualizing the clusters
plt.figure(figsize=(12, 6))
sns.scatterplot(x='Start_Lng', y='Start_Lat', hue='Cluster', data=new_data, palette='v')
plt.title('K-means Clustering (3 Clusters)')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()
```



```
In [39]: new_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2009085 entries, 0 to 2009084
Data columns (total 42 columns):
#   Column                                Dtype
---  -
0   ID                                    object
1   Source                               object
2   Severity                             int64
3   Start_Time                           object
4   End_Time                             object
5   Start_Lat                            float64
6   Start_Lng                            float64
7   Distance(mi)                         float64
8   Street                               object
9   City                                 object
10  County                               object
11  State                                object
12  Zipcode                             object
13  Country                             object
14  Timezone                             object
15  Airport_Code                         object
16  Weather_Timestamp                   object
17  Temperature(F)                      float64
18  Humidity(%)                         float64
19  Pressure(in)                        float64
20  Visibility(mi)                      float64
21  Wind_Direction                       object
22  Wind_Speed(mph)                     float64
23  Weather_Condition                    object
24  Amenity                              bool
25  Bump                                 bool
26  Crossing                             bool
27  Give_Way                             bool
28  Junction                             bool
29  No_Exit                             bool
30  Railway                             bool
31  Roundabout                          bool
32  Station                              bool
33  Stop                                 bool
34  Traffic_Calming                      bool
35  Traffic_Signal                       bool
36  Turning_Loop                         bool
37  Sunrise_Sunset                       object
38  Civil_Twilight                       object
39  Nautical_Twilight                    object
40  Astronomical_Twilight                 object
41  Cluster                              int32
dtypes: bool(13), float64(8), int32(1), int64(1), object(19)
memory usage: 461.8+ MB

```

```

In [40]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import silhouette_score

# Creating a smaller DataFrame with 5000 random rows
small_data = new_data.sample(n=5000, random_state=42)

# Extracting relevant columns for geographical clustering
geo_columns = ['Start_Lat', 'Start_Lng']
geo_data = small_data[geo_columns]

```

```

# Standardizing the data
scaler = StandardScaler()
geo_data_scaled = scaler.fit_transform(geo_data)

# Determining the optimal number of clusters using silhouette score and inertia
silhouette_scores_geo = []
inertia_values_geo = []

for n_clusters in range(2, 11):
    kmeans_geo = KMeans(n_clusters=n_clusters, random_state=42)
    cluster_labels_geo = kmeans_geo.fit_predict(geo_data_scaled)

    silhouette_avg_geo = silhouette_score(geo_data_scaled, cluster_labels_geo)
    silhouette_scores_geo.append(silhouette_avg_geo)

    inertia_value_geo = kmeans_geo.inertia_
    inertia_values_geo.append(inertia_value_geo)

# Finding the optimal number of clusters based on silhouette score
optimal_clusters_geo_silhouette = silhouette_scores_geo.index(max(silhouette_scores_geo))

# Find the optimal number of clusters based on inertia
optimal_clusters_geo_inertia = inertia_values_geo.index(min(inertia_values_geo)) + 2

# Plotting silhouette scores
plt.figure(figsize=(15, 6))

plt.subplot(1, 2, 1)
plt.plot(range(2, 11), silhouette_scores_geo, marker='o')
plt.title('Silhouette Scores for Different Number of Clusters')
plt.xlabel('Number of Clusters')
plt.ylabel('Silhouette Score')

# Plotting inertia values
plt.subplot(1, 2, 2)
plt.plot(range(2, 11), inertia_values_geo, marker='o', color='orange')
plt.title('Inertia Values for Different Number of Clusters')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')

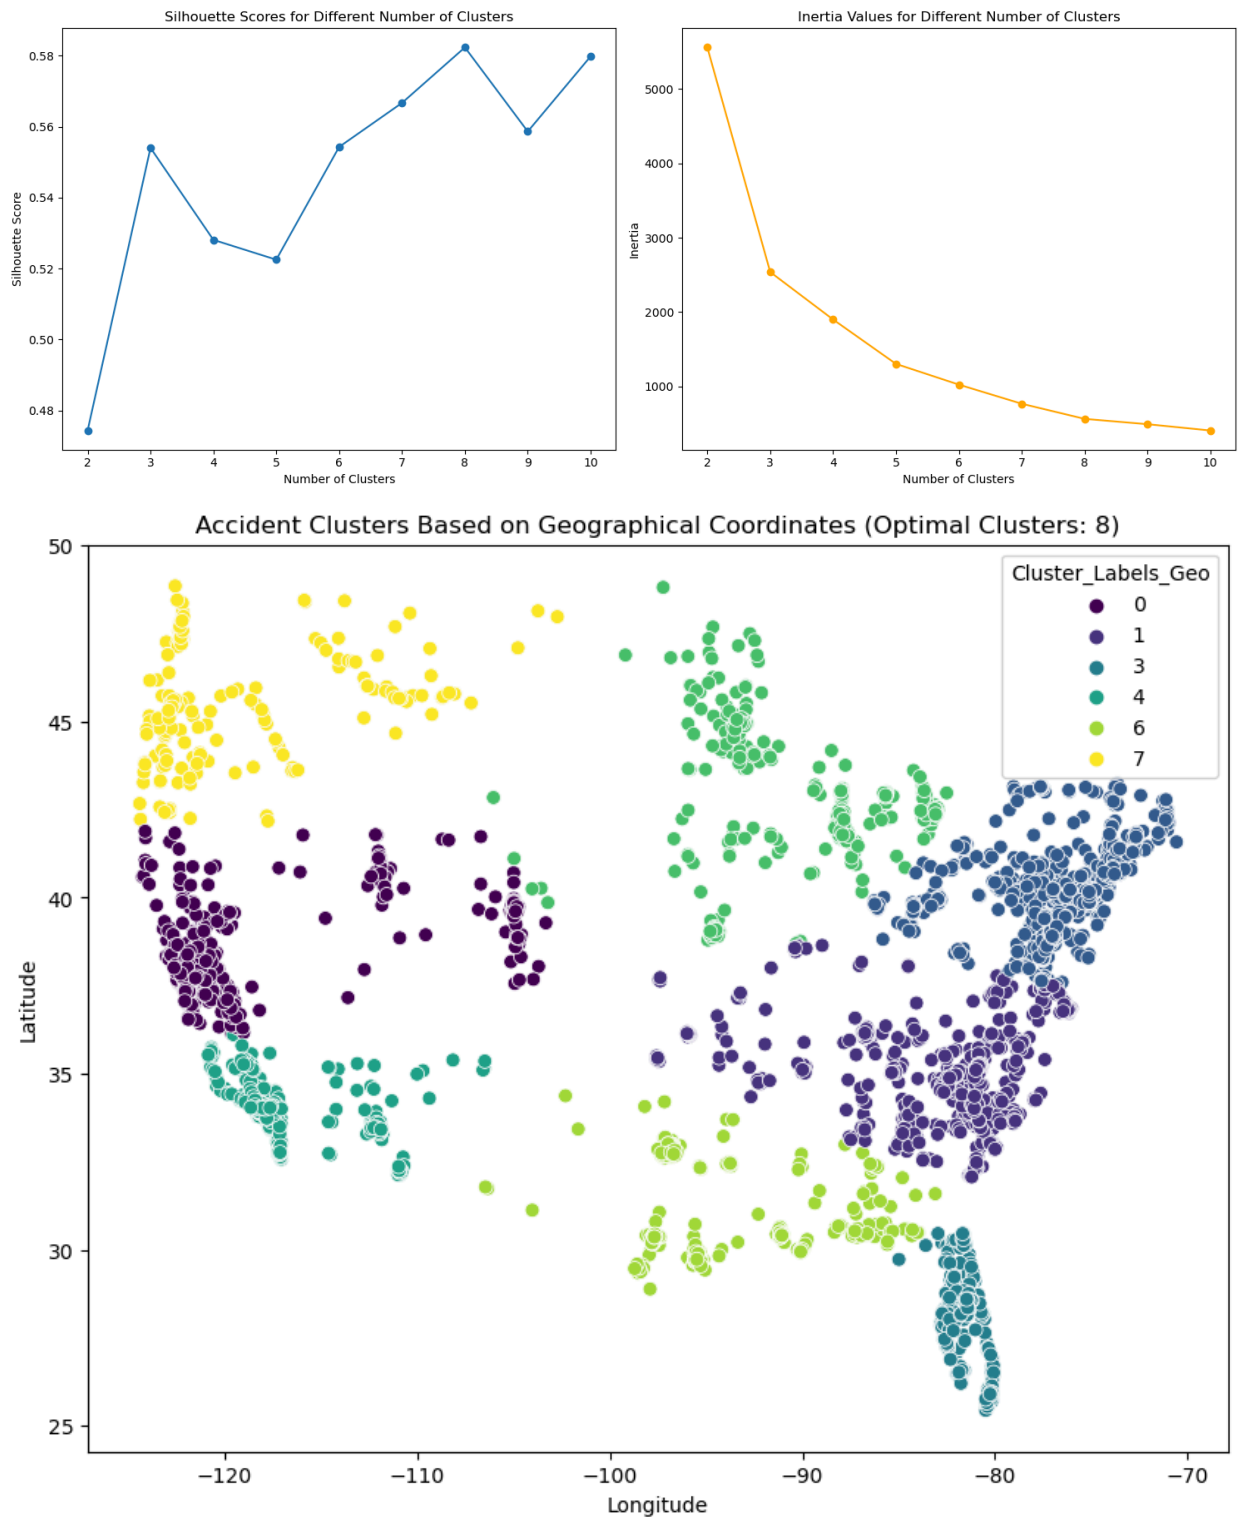
plt.tight_layout()
plt.show()

# Performing K-means clustering with the optimal number of clusters (We can choose either)
kmeans_geo_optimal = KMeans(n_clusters=optimal_clusters_geo_silhouette, random_state=42)
cluster_labels_geo_optimal = kmeans_geo_optimal.fit_predict(geo_data_scaled)

# Adding cluster labels to the original dataframe
small_data['Cluster_Labels_Geo'] = cluster_labels_geo_optimal

# Visualizing clusters on a scatter plot
plt.figure(figsize=(10, 8))
sns.scatterplot(x='Start_Lng', y='Start_Lat', hue='Cluster_Labels_Geo', data=small_data)
plt.title(f'Accident Clusters Based on Geographical Coordinates (Optimal Clusters: {optimal_clusters_geo_silhouette})')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()

```



```
In [41]: severity_column = 'Severity'

# 1. Cluster Statistics
cluster_statistics = small_data.groupby('Cluster_Labels_Geo')[severity_column].mean()
print("Cluster Statistics - Average Severity:")
print(cluster_statistics)

# 2. Spatial Patterns
plt.figure(figsize=(15, 8))
for cluster_label in range(optimal_clusters_geo_silhouette):
    cluster_data = small_data[small_data['Cluster_Labels_Geo'] == cluster_label]
    plt.scatter(cluster_data['Start_Lng'], cluster_data['Start_Lat'], label=f'Cluster
```

```

plt.title('Spatial Patterns Within Clusters (Silhouette Method)')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.legend()
plt.show()

# 3. Temporal Patterns
if 'Start_Time' in small_data.columns:
    small_data['Start_Time'] = pd.to_datetime(small_data['Start_Time'])
    small_data['Hour'] = small_data['Start_Time'].dt.hour

plt.figure(figsize=(15, 6))
sns.boxplot(x='Hour', y=severity_column, hue='Cluster_Labels_Geo', data=small_data)
plt.title('Temporal Patterns Within Clusters (Silhouette Method)')
plt.xlabel('Hour of the Day')
plt.ylabel('Severity')
plt.show()


```

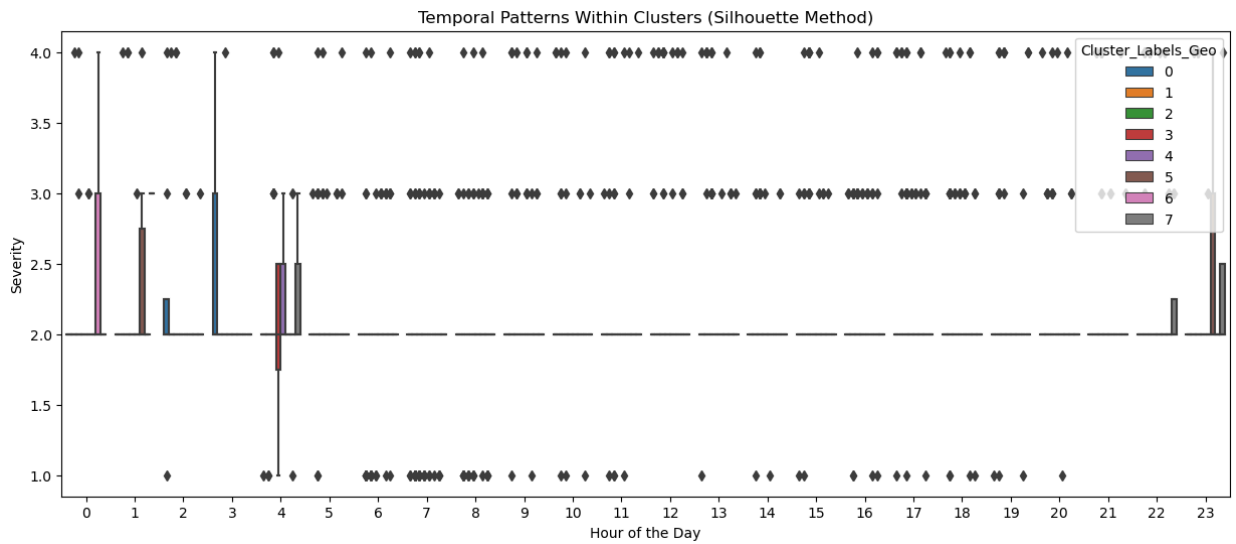
Cluster Statistics - Average Severity:

Cluster\_Labels\_Geo

0	2.043333
1	2.070696
2	2.121884
3	2.015267
4	2.051862
5	2.110837
6	2.076389
7	2.076923

Name: Severity, dtype: float64






```
In [42]: kmeans_geo_optimal = KMeans(n_clusters=optimal_clusters_geo_inertia , random_state=42)
cluster_labels_geo_optimal = kmeans_geo_optimal.fit_predict(geo_data_scaled)

# Adding cluster labels to the original dataframe
small_data['Cluster_Labels_Geo'] = cluster_labels_geo_optimal

# Visualizing clusters on a scatter plot
plt.figure(figsize=(10, 8))
sns.scatterplot(x='Start_Lng', y='Start_Lat', hue='Cluster_Labels_Geo', data=small_data)
plt.title(f'Accident Clusters Based on Geographical Coordinates (Optimal Clusters: {optimal_clusters_geo_inertia})')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()
```



```
In [43]: severity_column = 'Severity'

# 1. Cluster Statistics
cluster_statistics = small_data.groupby('Cluster_Labels_Geo')[severity_column].mean()
print("Cluster Statistics - Average Severity:")
print(cluster_statistics)

# 2. Spatial Patterns
plt.figure(figsize=(15, 8))
for cluster_label in range(optimal_clusters_geo_inertia):
    cluster_data = small_data[small_data['Cluster_Labels_Geo'] == cluster_label]
    plt.scatter(cluster_data['Start_Lng'], cluster_data['Start_Lat'], label=f'Cluster

plt.title('Spatial Patterns Within Clusters (WCSS inertia)')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.legend()
plt.show()

# 3. Temporal Patterns
if 'Start_Time' in small_data.columns:
    small_data['Start_Time'] = pd.to_datetime(small_data['Start_Time'])
    small_data['Hour'] = small_data['Start_Time'].dt.hour

plt.figure(figsize=(15, 6))
sns.boxplot(x='Hour', y=severity_column, hue='Cluster_Labels_Geo', data=small_data)
plt.title('Temporal Patterns Within Clusters (WCSS Inertia)')
```



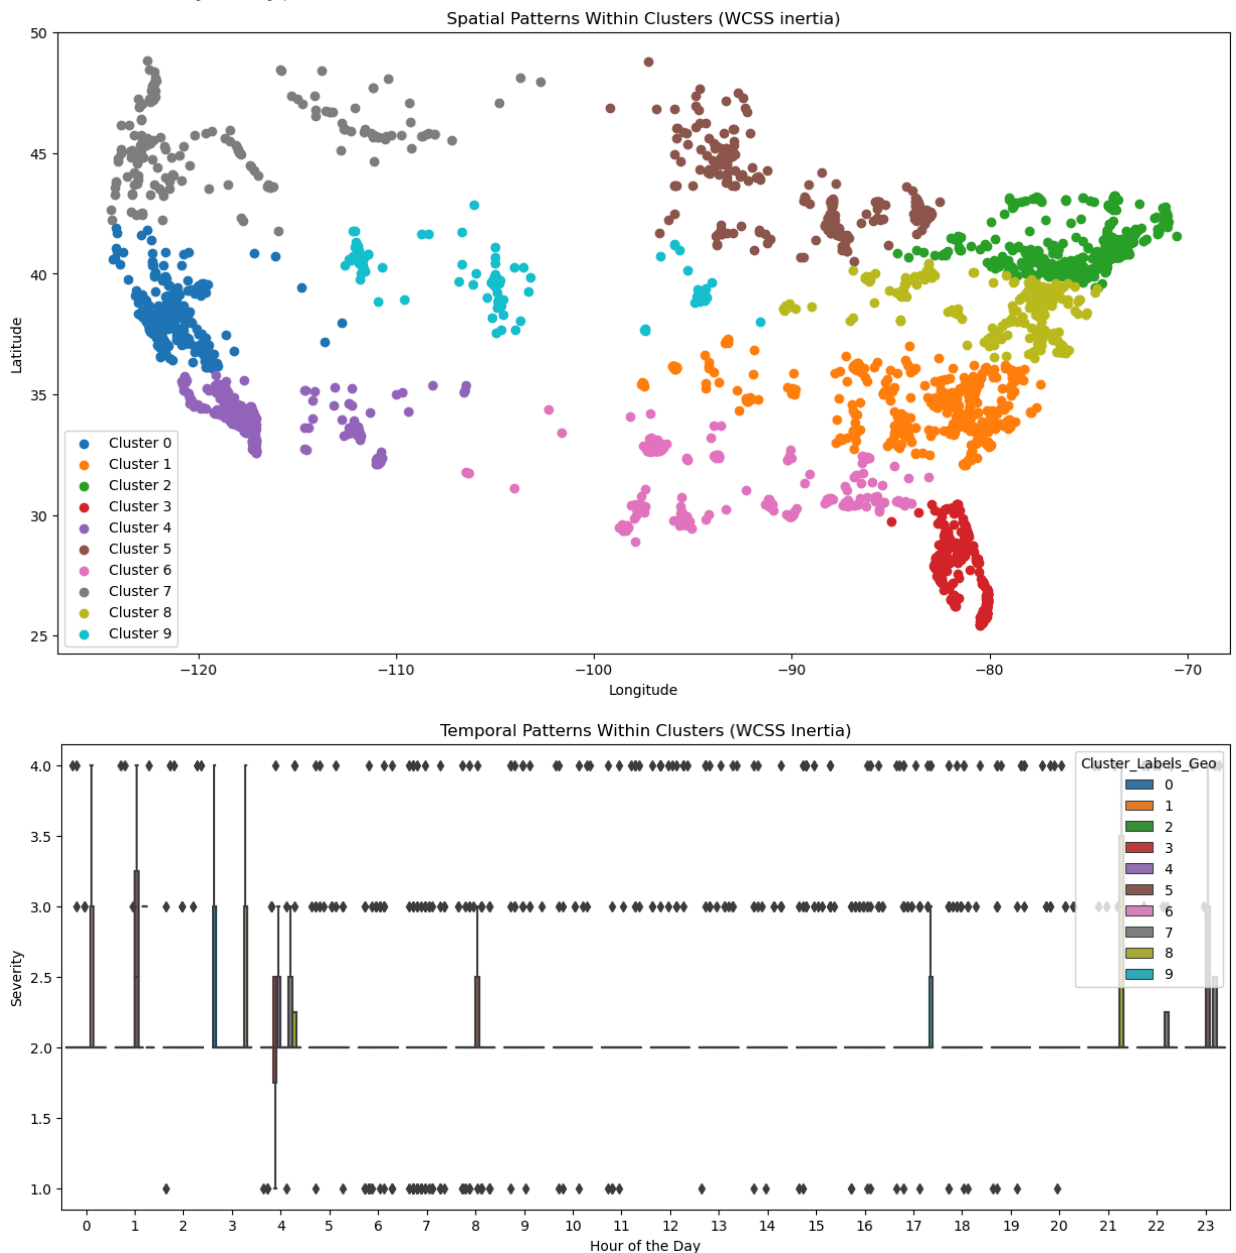
```
plt.xlabel('Hour of the Day')
plt.ylabel('Severity')
plt.show()
```

Cluster Statistics - Average Severity:

Cluster\_Labels\_Geo

```
0    2.026423
1    2.066757
2    2.093496
3    2.015267
4    2.052000
5    2.101744
6    2.076744
7    2.076531
8    2.155102
9    2.134503
```

Name: Severity, dtype: float64



```
In [49]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
```

```

from sklearn.metrics import silhouette_score
from scipy.cluster.hierarchy import linkage, dendrogram, fcluster
import matplotlib.pyplot as plt

selected_columns = ['Start_Lat', 'Start_Lng']

sample_size = 200
new_data_sample = new_data.sample(n=sample_size)

# Extracting the selected columns for clustering
X = new_data_sample[selected_columns]

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

cosine_dist = 1 - np.dot(X_scaled, X_scaled.T)

# Performing hierarchical clustering with ward linkage
linkage_matrix = linkage(cosine_dist, method='ward')

# Plotting the dendrogram
plt.figure(figsize=(15, 8))
dendrogram(linkage_matrix, labels=new_data_sample.index.values, orientation='top', dis
plt.title('Dendrogram of Hierarchical Clustering')
plt.xlabel('Data Points')
plt.show()

# Determining the optimal number of clusters based on the highest silhouette score
silhouette_scores = []
for k in range(2, 11):
    cluster_labels = fcluster(linkage_matrix, k, criterion='maxclust')
    silhouette_avg = silhouette_score(cosine_dist, cluster_labels)
    silhouette_scores.append(silhouette_avg)

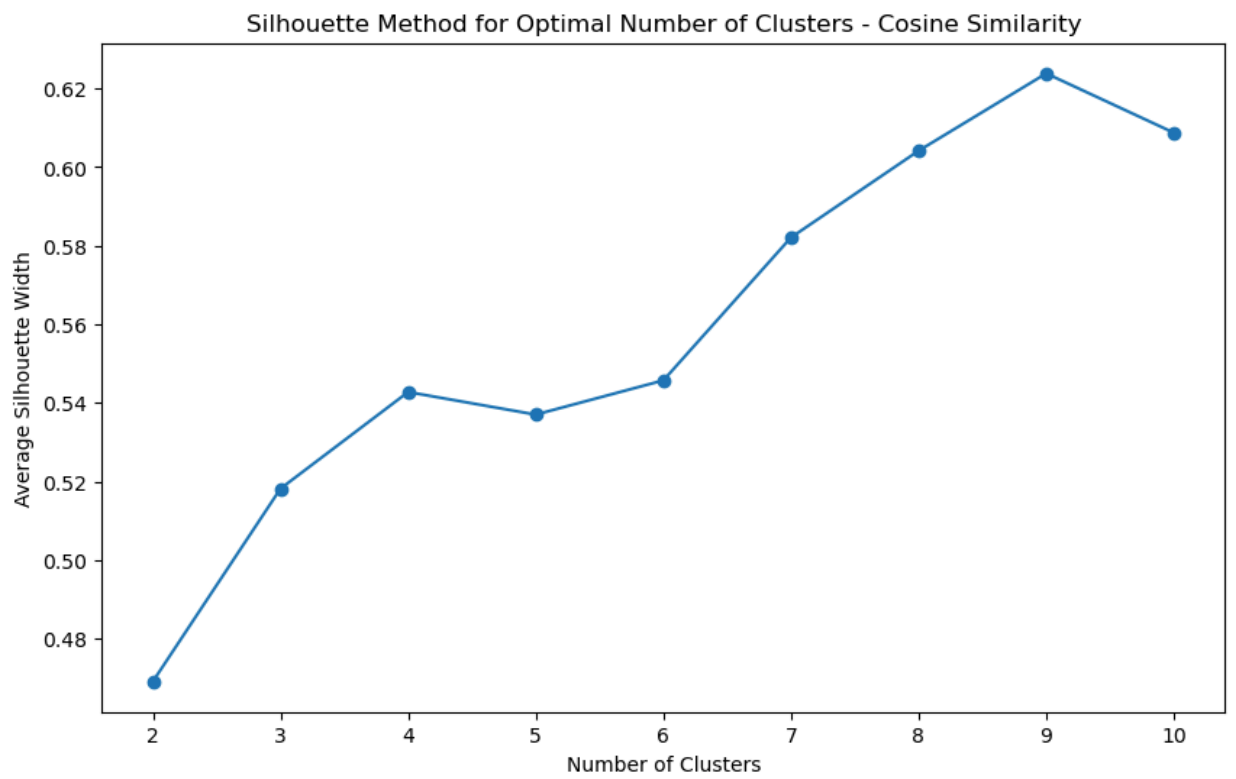

optimal_clusters = np.argmax(silhouette_scores) + 2

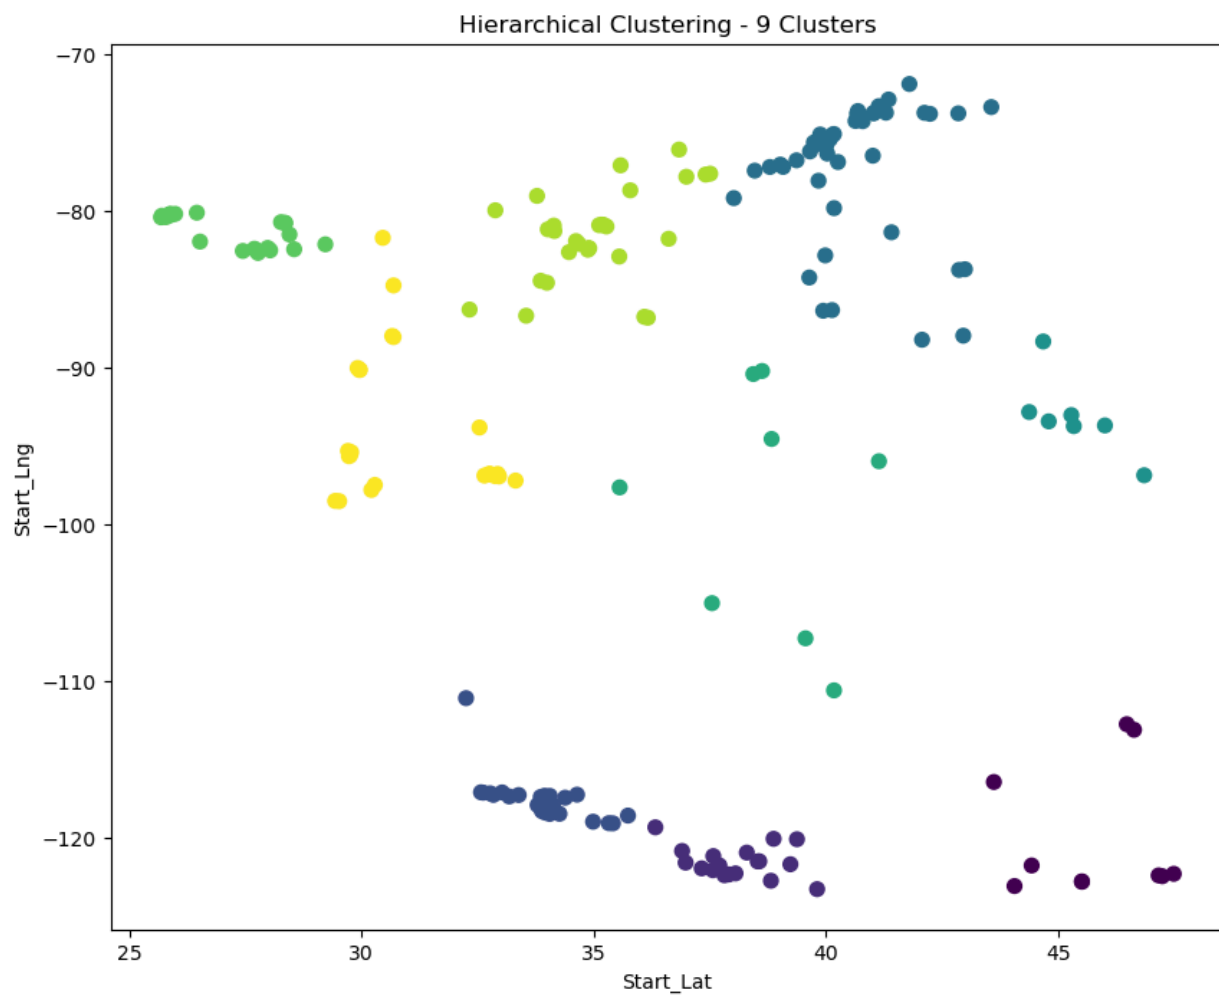
# Determining clusters
clusters = fcluster(linkage_matrix, optimal_clusters, criterion='maxclust')

# Plotting silhouette scores using a line plot
plt.figure(figsize=(10, 6))
plt.plot(range(2, 11), silhouette_scores, marker='o')
plt.title('Silhouette Method for Optimal Number of Clusters - Cosine Similarity')
plt.xlabel('Number of Clusters')
plt.ylabel('Average Silhouette Width')
plt.show()

# Visualizing clusters on a scatter plot
plt.figure(figsize=(10, 8))
plt.scatter(X.iloc[:, 0], X.iloc[:, 1], c=clusters, cmap='viridis', s=50)
plt.title(f'Hierarchical Clustering - {optimal_clusters} Clusters')
plt.xlabel('Start_Lat')
plt.ylabel('Start_Lng')
plt.show()

```





In [ ]: