

```
In [36]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

In [37]: df = pd.read_csv("last_two_years_accidents.csv")

In [38]: df.columns

Out[38]: Index(['ID', 'Source', 'Severity', 'Start_Time', 'End_Time', 'Start_Lat',
      'Start_Lng', 'Distance(mi)', 'Street', 'City', 'County', 'State',
      'Zipcode', 'Country', 'Timezone', 'Airport_Code', 'Weather_Timestamp',
      'Temperature(F)', 'Humidity(%)', 'Pressure(in)', 'Visibility(mi)',
      'Wind_Direction', 'Wind_Speed(mph)', 'Weather_Condition', 'Amenity',
      'Bump', 'Crossing', 'Give_Way', 'Junction', 'No_Exit', 'Railway',
      'Roundabout', 'Station', 'Stop', 'Traffic_Calming', 'Traffic_Signal',
      'Turning_Loop', 'Sunrise_Sunset', 'Civil_Twilight', 'Nautical_Twilight',
      'Astronomical_Twilight'],
      dtype='object')

In [39]: severity_labels = {
      1: 'Not Severe',
      2: 'Not Severe',
      3: 'Severe',
      4: 'Severe'
    }

In [40]: df['Severity'] = df['Severity'].map(severity_labels)

In [41]: size = df['Severity'].value_counts()['Severe']

In [42]: df_balanced_severity = pd.DataFrame()

In [43]: df_balanced_severity = df.groupby('Severity', group_keys = False).apply(lambda x: x.sample(size, random_state = 30))

In [44]: df_balanced_severity['Severity'].value_counts()

Out[44]: Severity
Not Severe    128292
Severe        128292
Name: count, dtype: int64

In [45]: categorical_features = set(["Weather_Condition", "Civil_Twilight", 'Wind_Direction'])

In [46]: for feature in categorical_features:
      df_balanced_severity[feature] = df_balanced_severity[feature].astype("category")

In [47]: bool_columns = df_balanced_severity.select_dtypes(include='bool').columns

In [48]: df_balanced_severity[bool_columns] = df_balanced_severity[bool_columns].replace({True:1, False:0})

In [49]: df2= df_balanced_severity[['Start_Lat','Start_Lng','Distance(mi)', 'Temperature(F)', 'Humidity(%)', 'Pressure(in)',
      'Visibility(mi)', 'Wind_Speed(mph)', 'Amenity', 'Bump', 'Crossing', 'Give_Way',
      'Junction', 'No_Exit', 'Railway', 'Roundabout', 'Station', 'Stop', 'Traffic_Calming', 'Traffic_Signal',
      'Civil_Twilight', 'Weather_Condition', 'Civil_Twilight',
      'Wind_Direction', 'Severity']]

In [50]: df2 = pd.get_dummies(df2, columns=list(categorical_features), drop_first=True)

In [51]: from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split

In [52]: Y = df2['Severity'] # target column
X = df2.drop(columns = ['Severity']) # features

In [53]: X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=30)

In [54]: sample_1 = df2.sample(5000, random_state=42)

In [55]: y_sample1 = sample_1["Severity"]
x_sample1 = sample_1.drop("Severity", axis=1)

In [56]: X_train, X_test, y_train, y_test = train_test_split(x_sample1, y_sample1, random_state=42)

In [57]: from sklearn import svm
```

Linear Kernel with Costs 0.1, 1, 10

```
In [58]: svc = svm.SVC(kernel= "linear", C= .2)
svc.fit(X_train, y_train)
```

```
Out[58]: SVC(C=0.2, kernel='linear')
```

```
In [59]: print("Train score:", svc.score(X_train, y_train))
print("Validation score:", svc.score(X_test, y_test))

Train score: 0.6170666666666667
Validation score: 0.5936
```


```
In [60]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, f1_score, roc_curve

y_pred = svc.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
Not Severe	0.63	0.51	0.56	639
Severe	0.57	0.69	0.62	611
accuracy			0.59	1250
macro avg	0.60	0.60	0.59	1250
weighted avg	0.60	0.59	0.59	1250

```
In [61]: y_pred = svc.predict(X_test)
confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)
```

```
In [62]: index = ["Not Severe Original", "Severe Original"]
columns = ["Not Severe Predicted", "Severe Predicted"]
conf_matrix = pd.DataFrame(data=confmat, columns=columns, index=index)
plt.figure(figsize=(4,4))
sns.heatmap(conf_matrix, annot=True, fmt="d")
plt.title("Confusion Matrix")
plt.show()
```



```
In [67]: svm_linear = svm.SVC(kernel= "linear", C= 1)
svm_linear.fit(X_train, y_train)

print("Train score:", svm_linear.score(X_train, y_train))
print("Validation score:", svm_linear.score(X_test, y_test))

y_pred = svm_linear.predict(X_test)


print(classification_report(y_test, y_pred))
```

Train score: 0.6176
Validation score: 0.592

	precision	recall	f1-score	support
Not Severe	0.63	0.50	0.56	639
Severe	0.57	0.68	0.62	611
accuracy			0.59	1250
macro avg	0.60	0.59	0.59	1250
weighted avg	0.60	0.59	0.59	1250

```
In [68]: y_pred = svm_linear.predict(X_test)
confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)
```

```
index = ["Not Severe Original", "Severe Original"]
columns = ["Not Severe Predicted", "Severe Predicted"]
conf_matrix = pd.DataFrame(data=confmat, columns=columns, index=index)
plt.figure(figsize=(4,4))
sns.heatmap(conf_matrix, annot=True, fmt="d")
plt.title("Confusion Matrix")
plt.show()
```



```
In [69]: svm_2 = svm.SVC(kernel= "linear", C= 10)
svm_2.fit(X_train, y_train)

print("Train score:", svm_2.score(X_train, y_train))
print("Validation score:", svm_2.score(X_test, y_test))

y_pred = svm_2.predict(X_test)


print(classification_report(y_test, y_pred))
```

Train score: 0.6242666666666666
Validation score: 0.5904

	precision	recall	f1-score	support
Not Severe	0.62	0.53	0.57	639
Severe	0.57	0.65	0.61	611
accuracy			0.59	1250
macro avg	0.59	0.59	0.59	1250
weighted avg	0.59	0.59	0.59	1250

```
In [70]: y_pred = svm_2.predict(X_test)
confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)

index = ["Not Severe Original", "Severe Original"]
columns = ["Not Severe Predicted", "Severe Predicted"]
conf_matrix = pd.DataFrame(data=confmat, columns=columns, index=index)
plt.figure(figsize=(4,4))
sns.heatmap(conf_matrix, annot=True, fmt="d")
plt.title("Confusion Matrix")
plt.show()
```



Polynomial Kernel With costs 0.1, 1, 10

```
In [71]: svm_poly = svm.SVC(kernel= "poly", C= 0.1, degree= 4)
svm_poly.fit(X_train, y_train)

print("Train score:", svm_poly.score(X_train, y_train))
print("Validation score:", svm_poly.score(X_test, y_test))

y_pred = svm_poly.predict(X_test)

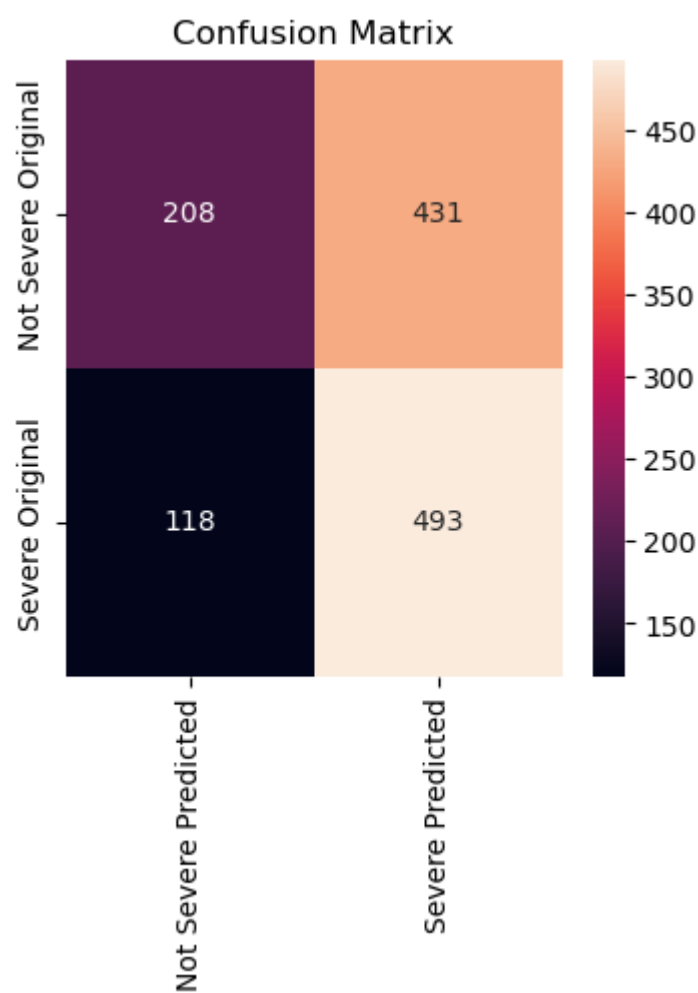
print(classification_report(y_test, y_pred))
```

```
Train score: 0.5664
Validation score: 0.5608
```

	precision	recall	f1-score	support
Not Severe	0.64	0.33	0.43	639
Severe	0.53	0.81	0.64	611
accuracy			0.56	1250
macro avg	0.59	0.57	0.54	1250
weighted avg	0.59	0.56	0.53	1250

```
In [72]: y_pred = svm_poly.predict(X_test)
confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)

index = ["Not Severe Original", "Severe Original"]
columns = ["Not Severe Predicted", "Severe Predicted"]
conf_matrix = pd.DataFrame(data=confmat, columns=columns, index=index)
plt.figure(figsize=(4,4))
sns.heatmap(conf_matrix, annot=True, fmt="d")
plt.title("Confusion Matrix")
plt.show()
```



```
In [75]: svm_poly_2 = svm.SVC(kernel= "poly", C= 1, degree= 4)
svm_poly_2.fit(X_train, y_train)

print("Train score:", svm_poly_2.score(X_train, y_train))
print("Validation score:", svm_poly_2.score(X_test, y_test))

y_pred = svm_poly_2.predict(X_test)

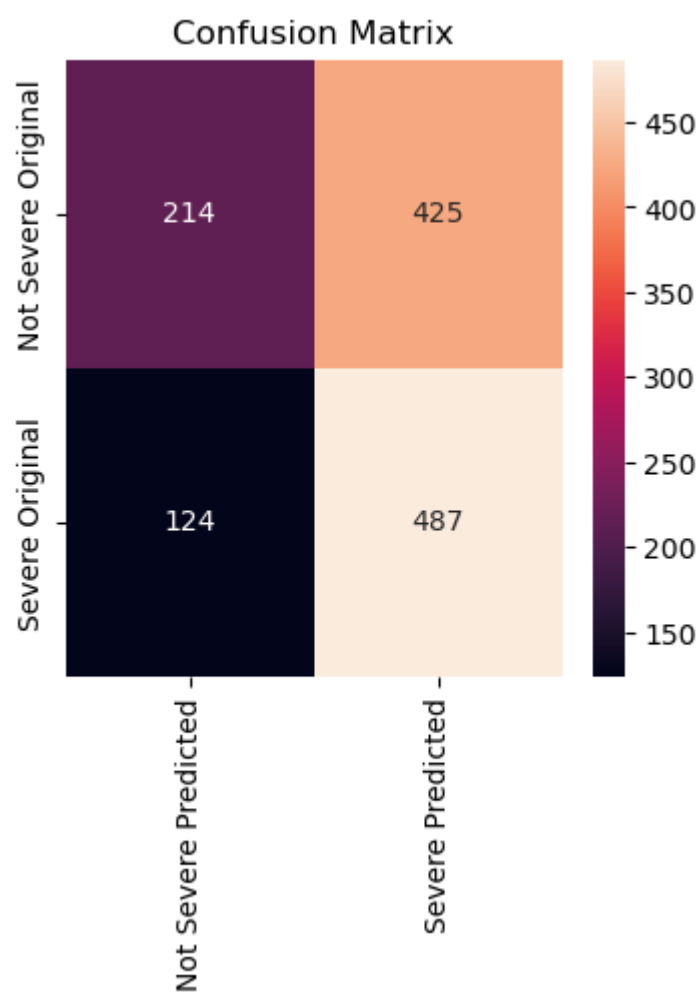
print(classification_report(y_test, y_pred))
```

Train score: 0.5690666666666667
Validation score: 0.5608

	precision	recall	f1-score	support
Not Severe	0.63	0.33	0.44	639
Severe	0.53	0.80	0.64	611
accuracy			0.56	1250
macro avg	0.58	0.57	0.54	1250
weighted avg	0.58	0.56	0.54	1250

```
In [74]: y_pred = svm_poly_2.predict(X_test)
confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)

index = ["Not Severe Original", "Severe Original"]
columns = ["Not Severe Predicted", "Severe Predicted"]
conf_matrix = pd.DataFrame(data=confmat, columns=columns, index=index)
plt.figure(figsize=(4,4))
sns.heatmap(conf_matrix, annot=True, fmt="d")
plt.title("Confusion Matrix")
plt.show()
```



```
In [76]: svm_poly_3 = svm.SVC(kernel= "poly", C= 10, degree= 4)
svm_poly_3.fit(X_train, y_train)

print("Train score:", svm_poly_3.score(X_train, y_train))
print("Validation score:", svm_poly_3.score(X_test, y_test))

y_pred = svm_poly_3.predict(X_test)


print(classification_report(y_test, y_pred))
```

Train score: 0.6056
Validation score: 0.596

	precision	recall	f1-score	support
Not Severe	0.64	0.48	0.55	639
Severe	0.57	0.72	0.63	611
accuracy			0.60	1250
macro avg	0.60	0.60	0.59	1250
weighted avg	0.61	0.60	0.59	1250

```
In [77]: y_pred = svm_poly_3.predict(X_test)
confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)

index = ["Not Severe Original", "Severe Original"]
columns = ["Not Severe Predicted", "Severe Predicted"]
conf_matrix = pd.DataFrame(data=confmat, columns=columns, index=index)
plt.figure(figsize=(4,4))
sns.heatmap(conf_matrix, annot=True, fmt="d")
plt.title("Confusion Matrix")
plt.show()
```



RBF Guassian Kernel With Costs 0.1, 1, 10

```
In [78]: svm_rbf = svm.SVC(kernel= "rbf", C= 0.1)
svm_rbf.fit(X_train, y_train)

print("Train score:", svm_rbf.score(X_train, y_train))
print("Validation score:", svm_rbf.score(X_test, y_test))

y_pred = svm_rbf.predict(X_test)

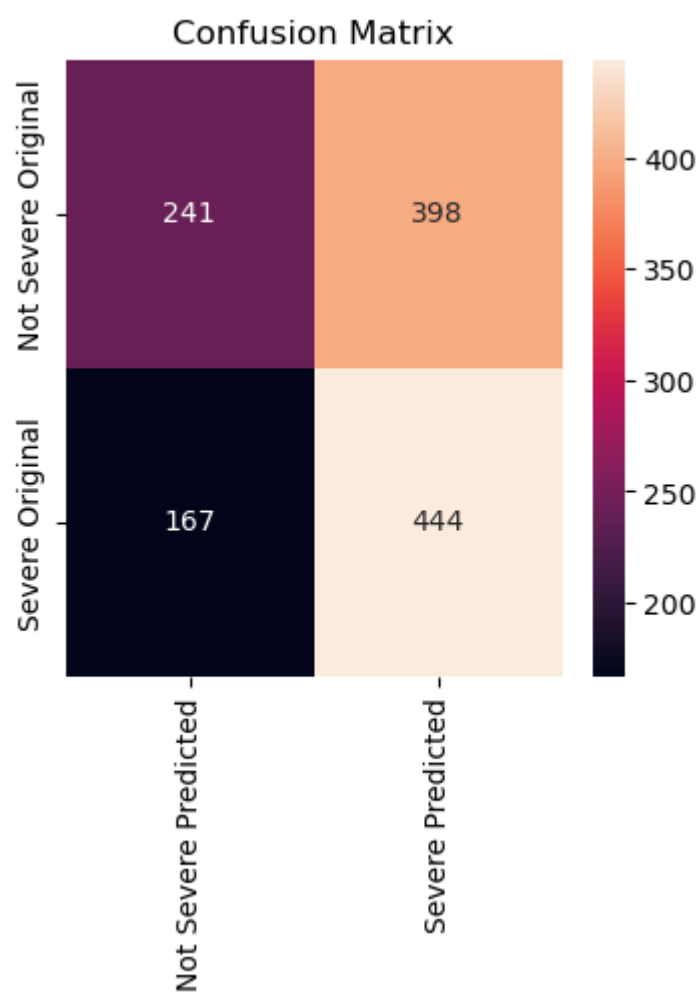
print(classification_report(y_test, y_pred))
```

```
Train score: 0.5421333333333334
Validation score: 0.548
```

	precision	recall	f1-score	support
Not Severe	0.59	0.38	0.46	639
Severe	0.53	0.73	0.61	611
accuracy			0.55	1250
macro avg	0.56	0.55	0.54	1250
weighted avg	0.56	0.55	0.53	1250

```
In [79]: y_pred = svm_rbf.predict(X_test)
confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)

index = ["Not Severe Original", "Severe Original"]
columns = ["Not Severe Predicted", "Severe Predicted"]
conf_matrix = pd.DataFrame(data=confmat, columns=columns, index=index)
plt.figure(figsize=(4,4))
sns.heatmap(conf_matrix, annot=True, fmt="d")
plt.title("Confusion Matrix")
plt.show()
```



```
In [80]: svm_rbf_2 = svm.SVC(kernel= "rbf", C= 1)
svm_rbf_2.fit(X_train, y_train)

print("Train score:", svm_rbf_2.score(X_train, y_train))
print("Validation score:", svm_rbf_2.score(X_test, y_test))

y_pred = svm_rbf_2.predict(X_test)

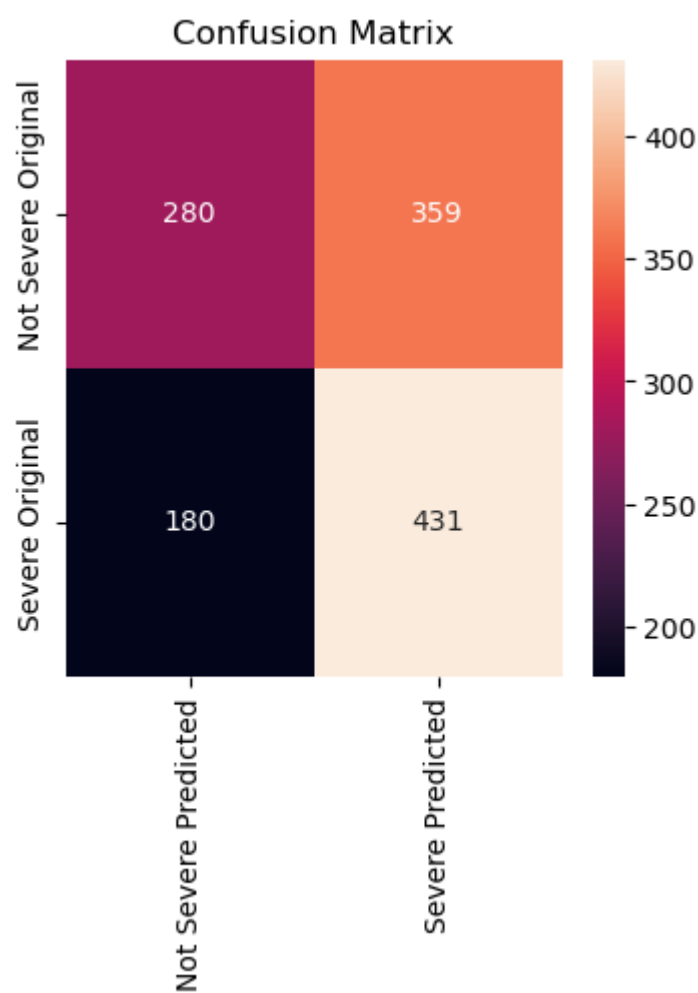
print(classification_report(y_test, y_pred))
```

Train score: 0.5813333333333334
Validation score: 0.5688

	precision	recall	f1-score	support
Not Severe	0.61	0.44	0.51	639
Severe	0.55	0.71	0.62	611
accuracy			0.57	1250
macro avg	0.58	0.57	0.56	1250
weighted avg	0.58	0.57	0.56	1250

```
In [81]: y_pred = svm_rbf_2.predict(X_test)
confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)

index = ["Not Severe Original", "Severe Original"]
columns = ["Not Severe Predicted", "Severe Predicted"]
conf_matrix = pd.DataFrame(data=confmat, columns=columns, index=index)
plt.figure(figsize=(4,4))
sns.heatmap(conf_matrix, annot=True, fmt="d")
plt.title("Confusion Matrix")
plt.show()
```

```
In [82]: sv_rbf_3 = svm.SVC(kernel= "rbf", C= 10)
sv_rbf_3.fit(X_train, y_train)

print("Train score:", sv_rbf_3.score(X_train, y_train))
print("Validation score:", sv_rbf_3.score(X_test, y_test))

y_pred = sv_rbf_3.predict(X_test)

print(classification_report(y_test, y_pred))
```

Train score: 0.6128
Validation score: 0.5896

	precision	recall	f1-score	support
Not Severe	0.63	0.47	0.54	639
Severe	0.56	0.71	0.63	611
accuracy			0.59	1250
macro avg	0.60	0.59	0.59	1250
weighted avg	0.60	0.59	0.58	1250

```
In [83]: y_pred = sv_rbf_3.predict(X_test)
confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)

index = ["Not Severe Original", "Severe Original"]
columns = ["Not Severe Predicted", "Severe Predicted"]
conf_matrix = pd.DataFrame(data=confmat, columns=columns, index=index)
plt.figure(figsize=(4,4))
sns.heatmap(conf_matrix, annot=True, fmt="d")
plt.title("Confusion Matrix")
plt.show()
```

